

▼ importing liabraries

```
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import matplotlib as plt
import matplotlib.pyplot as plt

cars = pd.read_csv("/content/auto_mpg.csv", names=None)

names = ['MPG', 'cylinders', 'Displacement', 'Horsepower', 'weight', 'Acceleration', 'Model Year', 'Origin', 'car Name']

cars.columns= names
```

```
cars.head()
```

	MPG	cylinders	Displacement	Horsepower	weight	Acceleration	Model Year	Origin	car Name
0	15.0	8	350.0	165	3693	11.5	70	1	buick skylark 320
1	18.0	8	318.0	150	3436	11.0	70	1	plymouth satellite
2	16.0	8	304.0	150	3433	12.0	70	1	amc rebel sst
3	17.0	8	302.0	140	3449	10.5	70	1	ford torino
4	15.0	8	429.0	198	4341	10.0	70	1	ford galaxie 500

Saved successfully!

(397, 9)

```
cars.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 397 entries, 0 to 396
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0    MPG             397 non-null   float64
1    cylinders        397 non-null   int64
2    Displacement     397 non-null   float64
3    Horsepower       397 non-null   object
4    weight           397 non-null   int64
5    Acceleration     397 non-null   float64
6    Model Year       397 non-null   int64
7    Origin           397 non-null   int64
8    car Name         397 non-null   object
dtypes: float64(3), int64(4), object(2)
memory usage: 28.0+ KB
```

```
cars["Horsepower"].isnull().sum()
```

0

```
cars=cars[cars["Horsepower"] != '?']
print(cars.shape)
```

(391, 9)

```
cars["Horsepower"]=cars["Horsepower"].astype('float')
print(cars.dtypes)
```

```
MPG             float64
cylinders        int64
Displacement     float64
Horsepower       float64
weight           int64
```

```

Acceleration    float64
Model Year      int64
Origin          int64
car Name        object
dtype: object
<ipython-input-10-05d0b84fe0b6>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view

```
cars["Horsepower"]=cars["Horsepower"].astype('float')
```

```
cars.isnull().sum()
```

```

MPG          0
cylinders    0
Displacement 0
Horsepower   0
weight       0
Acceleration 0
Model Year   0
Origin       0
car Name     0
dtype: int64

```

▼ Exploratory Analysis

```
cars["MPG"].describe()
```

```

count    391.000000
mean      23.459847
std        7.810128
min        9.000000
25%       17.000000

```

Saved successfully!

```

max      46.600000
Name: MPG, dtype: float64

```

```


import seaborn as sns
corrmatrix = cars.corr()
sns.heatmap(corrmatrix, square= True)

```

```

<ipython-input-13-c470deefec00>:2: FutureWarning: The default value of numeric_only in DataFrame.corr :
corrmatrix = cars.corr()
<Axes: >

```



```
from sklearn.metrics import mean_squared_error

features = ['cylinders','Displacement', 'Horsepower', 'weight','Acceleration','Model Year','Origin']
X = cars[features]
y = cars["MPG"]

X_train, X_test, y_train, y_test= train_test_split(X ,y, test_size=0.3,random_state=24)

regr =LinearRegression()
regr.fit(X_train, y_train)
```

▼ LinearRegression

LinearRegression()

```
predict_values= regr.predict(X_test)
```

predict_values

```
array([28.98319224, 26.39187916, 30.73892334, 17.54753575, 20.58585136,
       35.14751626, 29.07981513,  7.64073996, 25.54412284, 30.66116118,
       25.76382205, 10.311702  , 28.15407674, 20.68238756,  8.74694083,
       31.63534947, 31.16491326, 23.00018623, 24.09915332, 13.69535809,
       22.32753461, 24.58963406, 21.34618359, 28.64121716, 26.99022811,
       21.38332404, 26.62313065, 28.89539988, 36.13286973, 19.09332342,
       28.04784449, 30.59681937, 26.63209749, 11.65172323, 25.34278108,
       32.13890595, 31.32534439, 26.77727555, 12.32898894, 22.44821988,
       31.4342112  , 21.62424754, 26.05234347, 23.59296996, 28.74290616,
       17.07365939, 24.51134573, 23.47447188, 32.20856975, 25.69998714,
       19.00245374, 14.95690784, 26.39286337, 17.16613262, 25.62358536,
       30.31047022, 11.71617086, 32.66585635, 36.00997986, 26.58037419,
       19.2861861  , 24.19554281, 13.34353237, 23.06212512, 22.24769022,
       15.47359274, 22.72377863, 11.40965781, 29.13246907, 26.61513232,
       32.72317287, 24.88366239, 14.66364686, 29.98982358, 10.42772646,
       23.26801035, 16.52697645, 20.56773661, 22.08637274, 30.10965011,
       30.12040652, 20.05096527, 33.52760008, 30.08491344, 30.50880092,
       3.2459317393, 29.78036465, 27.96276329, 15.45911506, 31.01542712, 14.62852497,
       27.79287127, 17.27065529, 21.35730963, 27.80303003,  9.74876576,
       16.0608083  , 22.22404819,  6.45204073, 10.69077454, 23.85933183,
       29.5020639  , 22.71141174, 26.47561322, 33.46850582, 17.17617848,
       26.07545448, 32.27357293, 26.03072362, 26.91658339, 12.03018816,
       25.15404052, 10.0122611  , 20.95478654])
```

Saved successfully!

✕

```
pd.DataFrame(np.c_[X_test,y_test,predict_values],columns = ['cylinders','Displacement', 'Horsepower', 'weight','Acceleration','Model Year','Origin', 'MPG', 'mpg new'])
```

	cylinders	Displacement	Horsepower	weight	Acceleration	Model Year	Origin	MPG	mpg new
0	4.0	119.0	82.0	2720.0	19.4	82.0	1.0	31.0	28.983192
1	4.0	107.0	86.0	2464.0	15.5	76.0	2.0	28.0	26.391879
2	4.0	98.0	80.0	1915.0	14.4	79.0	1.0	35.7	30.738923
3	8.0	302.0	130.0	3870.0	15.0	76.0	1.0	13.0	17.547536
4	4.0	110.0	87.0	2672.0	17.5	70.0	2.0	25.0	20.585851
...
113	4.0	122.0	96.0	2300.0	15.5	77.0	1.0	25.5	26.916583
114	8.0	351.0	158.0	4363.0	13.0	73.0	1.0	13.0	12.030188
115	4.0	91.0	70.0	1955.0	20.5	71.0	1.0	26.0	25.154041
116	8.0	440.0	215.0	4312.0	8.5	70.0	1.0	14.0	10.012261
117	4.0	121.0	112.0	2868.0	15.5	73.0	2.0	19.0	20.954787

118 rows × 9 columns

```
from sklearn.metrics import r2_score
print(r2_score(y_test,predict_values))

0.8083852060985331

regr.score(X_test, y_test)
```

0.8083852060985331

▼ RandomForest

```
from sklearn.ensemble import RandomForestRegressor

forest_reg =RandomForestRegressor(random_state=12)
forest_reg.fit(X_train, y_train)
print("Training accuracy:",forest_reg.score(X_train,y_train))

y_pred= forest_reg.predict(X_test)
print("Testing accuracy:",forest_reg.score(X_test, y_test))

Training accuracy: 0.9818807392348246
Testing accuracy: 0.8943718899779411

MSE = mean_squared_error(y_test, predict_values)
print(MSE)

10.13592677140194
```

▼ Deployment

```
# save the Model
import joblib
joblib.dump(forest_reg, "Auto_MPG_predictor_model.pkl")
```

Saved successfully!

```
model= joblib.load("Auto_MPG_predictor_model.pkl")
```

```
model.predict(X_test)

array([27.916, 24.029, 32.798, 15.397, 23.081, 36.944, 35.688, 12.9 ,
       24.441, 28.135, 24.997, 13.395, 28.392, 20.019, 12.43 , 33.972,
       31.901, 24.246, 24.068, 14.375, 22.926, 25.052, 18.855, 29.479,
       25.967, 21.073, 27.225, 29.68 , 36.096, 18.663, 27.642, 35.618,
       24.609, 13.037, 22.63 , 35.748, 31.658, 29.175, 13.74 , 20.898,
       33.523, 17.63 , 22.778, 27.39 , 27.603, 18.022, 22.246, 22.519,
       31.336, 24.064, 18.921, 14.715, 24.699, 17.478, 27.787, 29.119,
       13.385, 33.3 , 36.399, 26.28 , 21.931, 20.942, 14.445, 22.335,
       21.263, 16.548, 19.582, 14.677, 27.61 , 23.728, 37.068, 26.125,
       14.33 , 30.995, 12.895, 19.552, 15.657, 18.308, 21.983, 31.454,
       17.91 , 28.662, 34.384, 31.701, 30.803, 32.855, 21.916, 21.778,
       24.23 , 30.21 , 18.695, 17.895, 14.27 , 32.024, 16.347, 29.998,
       16.648, 19.68 , 26.58 , 13.115, 18.031, 20.318, 12.75 , 14.122,
       24.11 , 31.629, 24.766, 25.259, 36.298, 15.903, 23.255, 33.745,
       24.041, 24.937, 13.915, 29.316, 14.55 , 21.683])
```

▼ launch & monitor

```
regr.save('/content/drive')
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-26-6059dd96a157> in <cell line: 1>()
----> 1 regr.save('/content/drive')

AttributeError: 'LinearRegression' object has no attribute 'save'
```

SEARCH STACK OVERFLOW

```
#path = '/content/drive//content/Auto_MPG_predictor_model.pkl'  
  
model.save(path)
```

▼ DEPLOYMENT

??

?

Saved successfully! ✕

! 0s completed at 7:24 PM

● ✕